

## ▼ Problem Statement 2 (Predict expected revenue from a customer)

The CRM team of a retail store wants to:

1. Understand which factors influence the revenue generated by the customers
2. Predict the revenue "Amount" for a given customer is likely to generate

```
# Importing libraries
```

```
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error
```

```
↳ /usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning:
    import pandas.util.testing as tm
```

```
# Reading the dataset
```

```
data=pd.read_csv("RetailCustomerRevenue.csv")
```

```
# Obtaining the column names
```

```
data.columns
```

```
↳ Index(['PersonID', 'Amount', 'FamilySize', 'Distance', 'Duration',
        'DirectVisits', 'OnlineVisits', 'Quantity', 'NumberofFrequentItems',
        'TransactionMode', 'Area', 'Occupation'],
        dtype='object')
```

```
# Reading the first few rows of data
```

```
data.head()
```

```
↳
```

	PersonID	Amount	FamilySize	Distance	Duration	DirectVisits	OnlineVisits	Quantity
0	C1104	3125	2	6	261	11	9	3
1	C1111	5298	2	5	323	9	9	2
2	C1117	4375	2	6	355	11	11	13
3	C1128	9700	5	7	418	51	41	29
4	C1132	3625	2	7	290	9	9	12

# Obtaining the dimensions

data.shape

↳ (2938, 12)

# Obtaining the data types

data.dtypes

↳

PersonID	object
Amount	int64
FamilySize	int64
Distance	int64
Duration	int64
DirectVisits	int64
OnlineVisits	int64
Quantity	int64
NumberOfFrequentItems	int64
TransactionMode	int64
Area	object
Occupation	int64
dtype:	object

# Summary Statistics

data.describe(include='all')

↳

	PersonID	Amount	FamilySize	Distance	Duration	DirectVisits	Onl:
<b>count</b>	2938	2938.000000	2938.000000	2938.000000	2938.000000	2938.000000	29
<b>unique</b>	2938	NaN	NaN	NaN	NaN	NaN	
<b>top</b>	C3960	NaN	NaN	NaN	NaN	NaN	
<b>freq</b>	1	NaN	NaN	NaN	NaN	NaN	

# Obtaining the variables having unique values less than 10

```
ins=[]
for a in data.columns:
    if(data[a].nunique(<10):
        ins.append(a)
        print(a,data[a].nunique())
        print(data[a].unique())
```

```
☐→ TransactionMode 2
    [2 1]
    Area 2
    ['Area1' 'Area2']
    Occupation 3
    [2 1 3]
```

ins

```
☐→ ['TransactionMode', 'Area', 'Occupation']
```

# Checking for missing values

```
data.isna().sum()
```

```
☐→ PersonID          0
    Amount            0
    FamilySize        0
    Distance          0
    Duration          0
    DirectVisits      0
    OnlineVisits      0
    Quantity          0
    NumberofFrequentItems  0
    TransactionMode   0
    Area              0
    Occupation        0
    dtype: int64
```

# Converting required data types into categorical

```
for col in ['TransactionMode', 'Area', 'Occupation', 'PersonID']:
    data[col]=data[col].astype('category')
```

```
# Checking to see if type has been converted
```

```
data.dtypes
```

```
↳ PersonID          category
   Amount            int64
   FamilySize        int64
   Distance           int64
   Duration           int64
   DirectVisits       int64
   OnlineVisits       int64
   Quantity           int64
   NumberofFrequentItems int64
   TransactionMode    category
   Area              category
   Occupation         category
   dtype: object
```

```
# Dropping duplicate values
```

```
data=data.drop_duplicates(keep='first')
```

```
# Dimensions of dataset
```

```
data.shape
```

```
↳ (2938, 12)
```

There are no duplicate values as the dimensions have not changed.

```
# Obtaining the correlation between each of the attributes
```

```
data.corr()
```

```
↳
```

Amount FamilySize Distance Duration DirectVisits OnlineVisits

# taking dependent variable "Amount" into y


```
y=np.array(data['Amount'])
```

```
Distance      -0.000649    0.436628    1.000000    -0.057219    -0.0004709    -0.0
```

### Obtaining Scatter plots between each attribute with y


```
plt.scatter(data['PersonID'],y)
```

☞ <matplotlib.collections.PathCollection at 0x7fe0219af3c8>



```
plt.scatter(data['FamilySize'],y)
```


☞ <matplotlib.collections.PathCollection at 0x7fe01fa3ecf8>



```
plt.scatter(data['Distance'],y)
```


☞

```
<matplotlib.collections.PathCollection at 0x7fe02006c748>
```




```
plt.scatter(data['Duration'],y)
```

```
↳ <matplotlib.collections.PathCollection at 0x7fe01fad8cc0>
```



```
plt.scatter(data['DirectVisits'],y)
```


```
↳ <matplotlib.collections.PathCollection at 0x7fe01fa10978>
```



```
plt.scatter(data['OnlineVisits'],y)
```


```
↳
```

```
<matplotlib.collections.PathCollection at 0x7fe020388a90>
```




```
plt.scatter(data['Quantity'],y)
```

```
↳ <matplotlib.collections.PathCollection at 0x7fe01f415588>
```



```
plt.scatter(data['TransactionMode'],y)
```

```
↳ <matplotlib.collections.PathCollection at 0x7fe020210fd0>
```



```
plt.scatter(data['Area'],y)
```


```
↳
```

<matplotlib.collections.PathCollection at 0x7fe020333c18>



```
plt.scatter(data['Occupation'],y)
```

↳ <matplotlib.collections.PathCollection at 0x7fe01fa385c0>



According to the scatter plots, "TransactionMode", "Area" and "Occupation" are categorical.

```
# Splitting into train and test data
```

```
train_data,test_data=train_test_split(data,test_size=0.2,random_state=123)
```

```
# Printing the dimensions of the train and test data
```

```
print(train_data.shape)
```

```
print(test_data.shape)
```

↳ (2350, 12)  
(588, 12)

```
# Creating dummy variables
```

```
train_dum=pd.get_dummies(columns=['TransactionMode','Area', 'Occupation'],data=train_data)
```

```
test_dum=pd.get_dummies(columns=['TransactionMode','Area', 'Occupation'],data=test_data)
```



```
# Printing the columns of the data after creating the dummy variables
```

```
print(train_dum.columns)
print(test_dum.columns)
```

```
↳ Index(['PersonID', 'Amount', 'FamilySize', 'Distance', 'Duration',
        'DirectVisits', 'OnlineVisits', 'Quantity', 'NumberofFrequentItems',
        'TransactionMode_1', 'TransactionMode_2', 'Area_Area1', 'Area_Area2',
        'Occupation_1', 'Occupation_2', 'Occupation_3'],
        dtype='object')
Index(['PersonID', 'Amount', 'FamilySize', 'Distance', 'Duration',
        'DirectVisits', 'OnlineVisits', 'Quantity', 'NumberofFrequentItems',
        'TransactionMode_1', 'TransactionMode_2', 'Area_Area1', 'Area_Area2',
        'Occupation_1', 'Occupation_2', 'Occupation_3'],
        dtype='object')
```

```
train_data1=train_dum
test_data1=test_dum
```

```
# Dropping "PersonID" attribute
```

```
train_data1=train_dum.drop('PersonID',axis=1)
test_data1=test_dum.drop('PersonID',axis=1)
```

```
print(train_data1.columns)
print(test_data1.columns)
```

```
↳ Index(['Amount', 'FamilySize', 'Distance', 'Duration', 'DirectVisits',
        'OnlineVisits', 'Quantity', 'NumberofFrequentItems',
        'TransactionMode_1', 'TransactionMode_2', 'Area_Area1', 'Area_Area2',
        'Occupation_1', 'Occupation_2', 'Occupation_3'],
        dtype='object')
Index(['Amount', 'FamilySize', 'Distance', 'Duration', 'DirectVisits',
        'OnlineVisits', 'Quantity', 'NumberofFrequentItems',
        'TransactionMode_1', 'TransactionMode_2', 'Area_Area1', 'Area_Area2',
        'Occupation_1', 'Occupation_2', 'Occupation_3'],
        dtype='object')
```

```
# Placing the independent variables into ind_atr
```

```
ind_atr=list(set(train_data1.columns)-set(['Amount']))
```

```
ind_atr
```

```
↳
```

```

    ['Occupation_3',
     'Duration',
     'OnlineVisits',
     'Area_Area1',
     'FamilySize',
     'Area_Area2',
     'DirectVisits',
     'Occupation_1',
     'TransactionMode_1',
     'Quantity',
     'Occupation 2',

# Fitting LinearRegression model

linreg=LinearRegression()
res_sklearn=linreg.fit(train_data1[ind_atr],train_data1['Amount'])

# Predicting train and test values from the model

pres_train_skleran=res_sklearn.predict(train_data1[ind_atr])
pres_test_skleran=res_sklearn.predict(test_data1[ind_atr])

# Printing the coefficients

res_sklearn.coef_

↳ array([-1.37942112e+01, -5.25526037e-01,  2.32211345e+02, -1.51619314e+02,
          4.71029449e+01,  1.51619314e+02,  2.62546923e+02, -1.82477056e+00,
          -1.32104193e+02,  3.84743602e-02,  1.56189818e+01, -9.47350161e+00,
          -2.91464118e+02,  1.32104193e+02])

# Finding mean absolute error for train data

mean_absolute_error(np.array(train_data1['Amount']),pres_train_skleran)

↳ 793.2396053830885

# Finding mean absolute error for test data

mean_absolute_error(np.array(test_data1['Amount']),pres_test_skleran)

↳ 778.8292727155123

Amount depends on 'FamilySize', 'Distance', 'Duration', 'DirectVisits', 'OnlineVisits', 'Quantity',
'NumberofFrequentItems', 'TransactionMode', 'Area', 'Occupation'

```

## Stats model

```
ind_atr1=list(set(data.columns)-set(['Amount','PersonID']))
```

```
ind_atr1
```

```
↳ ['Occupation',  
    'Duration',  
    'OnlineVisits',  
    'FamilySize',  
    'Area',  
    'DirectVisits',  
    'TransactionMode',  
    'Quantity',  
    'Distance',  
    'NumberofFrequentItems']
```

```
#Creating the formula
```

```
x='+'.join(ind_atr1)  
x
```

```
↳ 'Occupation+Duration+OnlineVisits+FamilySize+Area+DirectVisits+TransactionMode+Quantit  
v+Distance+NumberofFrequentItems'
```

```
formula="~".join(('Amount',x))  
print(formula)
```

```
↳ Amount~Occupation+Duration+OnlineVisits+FamilySize+Area+DirectVisits+TransactionMode+Qu
```

```
from statsmodels.formula.api import ols
```

```
#Model fit and summary
```

```
lm_mod=ols(formula=formula,data=data)  
result=lm_mod.fit()
```

```
print(result.summary2())
```

```
↳
```

Results: Ordinary least squares						
=====						
Model:	OLS	Adj. R-squared:	0.716			
Dependent Variable:	Amount	AIC:	49518.8936			
Date:	2020-07-15 11:09	BIC:	49590.7194			
No. Observations:	2938	Log-Likelihood:	-24747.			
Df Model:	11	F-statistic:	673.3			
Df Residuals:	2926	Prob (F-statistic):	0.00			
R-squared:	0.717	Scale:	1.2179e+06			
-----						
	Coef.	Std.Err.	t	P> t	[0.025	0.975]
-----						
Intercept	621.3437	120.6836	5.1485	0.0000	384.7103	857.9771
Occupation[T.2]	4.6484	50.9386	0.0913	0.9273	-95.2308	104.5276
Occupation[T.3]	-13.7002	49.8891	-0.2746	0.7836	-111.5215	84.1212
Area[T.Area2]	267.0128	66.3723	4.0230	0.0001	136.8716	397.1540
TransactionMode[T.2]	299.4014	55.5814	5.3867	0.0000	190.4188	408.3841
Duration	-0.5462	0.2379	-2.2961	0.0217	-1.0126	-0.0798
OnlineVisits	232.7890	8.3187	27.9837	0.0000	216.4778	249.1001
Condition Number	21.7605	22.4214	4.4465	0.1567	12.0000	75.7000

OBSERVATION

Distance	-1.1675	7.6506	-0.5117	0.5860	-19.1686	10.8335
----------	---------	--------	---------	--------	----------	---------

The R-Squared error is 0.717. If R-squared value  $r > 0.7$  this value is generally considered strong effect size.

Prob(>0.001005):	0.000	Prob(>0.001005):	0.000
1.000	1.000	1.000	1.000

-----  
\* The condition number is large (2e+04). This might indicate strong multicollinearity or other numerical problems.